

CS539 - NLP - Homework 2

Nuttaree Busarapongpanich

1 Demystifying Recurrent Neural Networks

1.1 Deriving LSTM Gradients for Univariate Case

TASK 1.1

The gradient for h_t w.r.t. c_t :

$$\frac{\delta h_t}{\delta c_t} = o_t * (1 - \tanh^2(c_t)) \quad (1)$$

The gradient for c_t w.r.t. i_t :

$$\frac{\delta c_t}{\delta i_t} = g_t \quad (2)$$

The gradient for a i_t w.r.t. w_{ix} :

$$\frac{\delta i_t}{\delta w_{ix}} = \sigma(w_{ix}x_t + w_{ih}h_{t-1} + b_i) * (1 - \sigma(w_{ix}x_t + w_{ih}h_{t-1} + b_i)) * x_t \quad (3)$$

The gradient for a i_t w.r.t. b_i :

$$\frac{\delta i_t}{\delta b_i} = \sigma(w_{ix}x_t + w_{ih}h_{t-1} + b_i) * (1 - \sigma(w_{ix}x_t + w_{ih}h_{t-1} + b_i)) \quad (4)$$

The gradient for h_t w.r.t. w_{ix} . According to chain rule:

$$\frac{\delta h_t}{\delta w_{ix}} = \frac{\delta h_t}{\delta c_t} * \frac{\delta c_t}{\delta i_t} * \frac{\delta i_t}{\delta w_{ix}} \quad (5)$$

Substitute (5) by (1), (2), (3):

$$\frac{\delta h_t}{\delta w_{ix}} = o_t * (1 - \tanh^2(c_t)) * g_t * \sigma(w_{ix}x_t + w_{ih}h_{t-1} + b_i) * (1 - \sigma(w_{ix}x_t + w_{ih}h_{t-1} + b_i)) * x_t \quad (6)$$

Next, we want to find the gradient for h_t w.r.t. b_i . According to chain rule:

$$\frac{\delta h_t}{\delta b_i} = \frac{\delta h_t}{\delta c_t} * \frac{\delta c_t}{\delta i_t} * \frac{\delta i_t}{\delta b_i} \quad (7)$$

Substitute (7) by (1), (2), (4):

$$\frac{\delta h_t}{\delta b_i} = o_t * (1 - \tanh^2(c_t)) * g_t * \sigma(w_{ix}x_t + w_{ih}h_{t-1} + b_i) * (1 - \sigma(w_{ix}x_t + w_{ih}h_{t-1} + b_i)) \quad (8)$$

TASK 1.1 extra credit

The gradient for a loss w.r.t. i_t

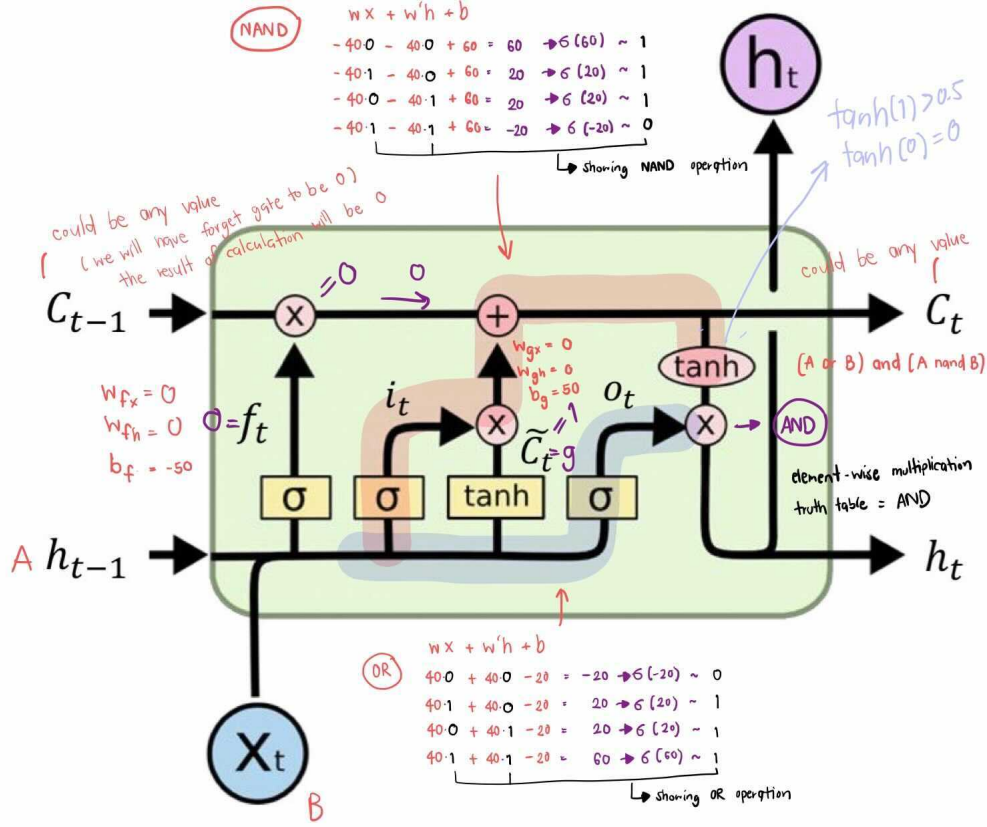


Figure 1: The flow with weights and biases of the Network for Parity

2 Learning to Copy Finite State Machines

2.1 Parity: Examining Generalization To Longer Sequences

TASK 2.1

In the implementation of a forward function for LSTM model for the parity task, there are 2 inputs for the forward function which is x (a padded input) and s (list of sequent lengths). The forward function contains 3 parts: a `pack_padded_sequence` function, a standard LSTM, a linear layer.

The input that we get as a parameter input to the forward function has a different dimension from what the `pack_padded_sequence` function requires, so we have to `unsqueeze` the input x to be able to parse to the `pack_padded_sequence` function without an error.

After the input dimension was updated, I then call `pack_padded_sequence` to optimize the computations, so we don't waste time and resources to compute unnecessary value. Then the output of this function will be the input of a standard LSTM function. I am using a built-in PyTorch LSTM function with the setting: `batch_first = True`, `input_size = 1`, `hidden_size = hidden_dim`, `num_layers = 1`. One of the output from this function is the hidden state, which I use the last index of the hidden state to the linear layer. The linear function has the setting: `in_features = hidden_dim`, `out_features = 2`. The output of the linear function is the output of this forward function.

We can see that I did not use the unpacking function for this task. The parity task is a sentence-level classification, so we do not care about the immediate hidden states.

With a standard LSTM and a linear layer, the program yields the result as below. We can see that the train accuracy reaches to 100% since around 160 epochs.

1	2021-02-01	08:25:53	INFO	epoch 130	train loss 0.393, train acc 0.855
2	2021-02-01	08:25:53	INFO	epoch 140	train loss 0.263, train acc 0.935
3	2021-02-01	08:25:53	INFO	epoch 150	train loss 0.150, train acc 0.984
4	2021-02-01	08:25:53	INFO	epoch 160	train loss 0.078, train acc 1.000
5	2021-02-01	08:25:53	INFO	epoch 170	train loss 0.040, train acc 1.000
6	2021-02-01	08:25:53	INFO	epoch 180	train loss 0.022, train acc 1.000
7	2021-02-01	08:25:53	INFO	epoch 190	train loss 0.014, train acc 1.000
8	2021-02-01	08:25:53	INFO	epoch 200	train loss 0.010, train acc 1.000
9	2021-02-01	08:25:53	INFO	epoch 210	train loss 0.008, train acc 1.000

TASK 2.2




Figure 2: The validation accuracy rate when `hidden_dim = 64`, `epochs = 2000`, `rl = 0.003`




Figure 3: The validation accuracy rate when `hidden_dim = 32`, `epochs = 2000`, `rl = 0.003`

We trained the model on `max_length` equals to 5. The performance on the validation set is also very good when the `max_length` is around 5. When the length comes to more than 15, the model seems to lose the generalization (overfit problem). When I reduce the hidden state size to 32, the validation accuracy rate is 100% until binary sequence of length 20.

TASK 2.3

From Figure 4-9, we can see that when the hidden dimension is smaller, the learning rate has to be higher to be able to maintain the high validation accuracy. It is easy to compare in the Figure 7 and 8, where the hidden state dimension is 2 with the epochs equal to 2000 in the different learning rate. The learning rate 0.05 yields the better validation accuracy rate than 0.03. I can reduce the hidden state dimension to 1 (The smallest size I can get) but the learning rate will be 0.8. When I reduce dimension, the model tends to learn slower, which would be a reason why we need to increase the learning rate.




Figure 4: The validation accuracy rate when $\text{hidden_dim} = 200$, $\text{epochs} = 2000$, $\text{rl} = 0.003$




Figure 5: The validation accuracy rate when $\text{hidden_dim} = 80$, $\text{epochs} = 2000$, $\text{rl} = 0.003$




Figure 6: The validation accuracy rate when $\text{hidden_dim} = 10$, $\text{epochs} = 2000$, $\text{rl} = 0.03$




Figure 7: The validation accuracy rate when $\text{hidden_dim} = 5$, $\text{epochs} = 2000$, $\text{rl} = 0.03$




Figure 8: The validation accuracy rate when $\text{hidden_dim} = 3$, $\text{epochs} = 2000$, $\text{rl} = 0.03$




Figure 9: The validation accuracy rate when $\text{hidden_dim} = 2$, $\text{epochs} = 2000$, $\text{rl} = 0.03$

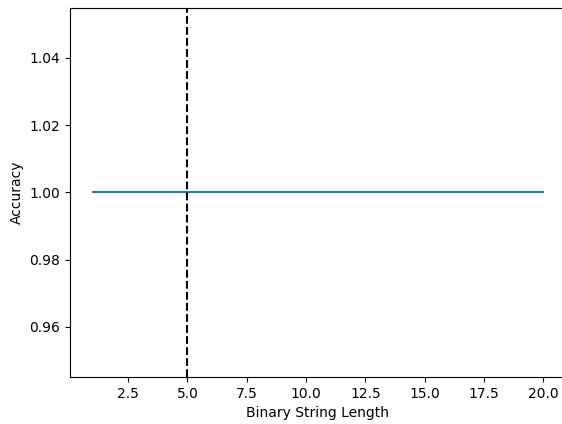


Figure 10: The validation accuracy rate when `hidden_dim = 2`, `epochs = 2000`, `rl = 0.05`




Figure 11: The validation accuracy rate when `hidden_dim = 1`, `epochs = 2000`, `rl = 0.8`

TASK 2.4

Vanilla RNNs have a serious problem in vanish gradient. The cause of this problem is vanilla RNNs calculate the gradient through all layers and through time. When vanilla RNNs was trained on the long distance and the time past, they could not use some of the earlier information. On the other hand, LSTMs has the cell states can be used as a bypass through the network. This is the reason that LSTMs has the capability to remember the longer sequences than vanilla RNNs.

When we evaluate the validation accuracy, the vanilla RNN couldn't remember the earlier state of the model so it seems to train base on the later information. As the model tries to fit the training set, this is why it impact only the validation set.

3 Part-of-Speech Tagging

TASK 3.1

The first 1,000 sentences' topic is related to the politic. For example:

pakistan now has a large force deployed in baluchistan which was not there before.
afghanistan has had the benefit of cooperation from both the us and iran
he clearly enjoyed , as governor , watching executions
this year a newspaper claimed two uk firms were involved in a deal in which thousands of guns for iraqi forces were re-routed to al - qaeda

For the 2,000th-3,000th sentence, the format is the email format. For example:

mary.ellenberger@enron.com on 05/03/2001 04:06:52 pm please respond to mary.ellenberger@enron.com cc : subject : re : if tgpl la z1 jan 9.95 feb 6.25 mar 4.98 april 5.37 esimien@nisource.com on 05/03/2001 01:32:35 pm hi mary , please do me a favour and give me the subject price for jan , feb , mar and apr 2001 . thanks , gregg penman
note : the information in this email is confidential and may be legally privileged .

For the 9,000th-10,000th sentence, They are reviews. For example:

because of the ants i dropped them to a 3 star .
cj and company did all that we ask and 10 times more .
excellent customer service and honest feedback .

These information, that stay in groups by the order of examples, will give us an idea why we should shuffle the data. If we do not shuffle the data, the model will learn group by group and eventually the model will have bad performance.




Figure 12: The histogram of part-of-speech tags

1	Tag	Count	Percentage
2			
3	NOUN	34781	17.00 %
4	PUNCT	23679	11.57 %
5	VERB	23081	11.28 %
6	PRON	18577	9.08 %
7	ADP	17638	8.62 %
8	DET	16285	7.96 %
9	PROPN	12946	6.33 %
10	ADJ	12477	6.10 %
11	AUX	12343	6.03 %
12	ADV	10548	5.16 %
13	CCONJ	6707	3.28 %
14	PART	5567	2.72 %
15	NUM	3999	1.95 %
16	SCONJ	3843	1.88 %
17	X	847	0.41 %
18	INTJ	688	0.34 %
19	SYM	599	0.29 %

We can see from the histogram in Figure 12 and the result of counting number of tags that the data is not balance. The top 3 highest amount of tags are noun, verb, and punctuation, while interjection, symbol, and other types are 3 smallest in size.

Lemmatization and stemming are not appropriate approaches for this dataset. They will make the model even worse. For example, lemmatization could change the adjective form to noun form, and the model will lose the information of predicting either noun or adjective. Stemming will cut the end part of the word, which we will lost the information of part-of-speech tagging. We could normalize but it is not necessary. For example, when we normalize from n't to not, after model learn, they might give n't or not to the same tag.

It is obvious that text preprocessing depends on both dataset and task/problem.

For this dataset, the tag that has highest percentage is NOUN, which is 17%. If the model always return NOUN for every word, we will have the word-level accuracy to be 17%.

TASK 3.2

The train/validation/test accuracies from `driver_udpos.py` are shown as below. The test accuracy rate is around 89.356%.

1	epoch 1	train loss 1.541, train acc 57.305%	val loss 0.902, val acc 77.084%
2	epoch 2	train loss 0.450, train acc 88.069%	val loss 0.554, val acc 85.500%
3	epoch 3	train loss 0.268, train acc 92.611%	val loss 0.465, val acc 87.517%
4	epoch 4	train loss 0.202, train acc 94.294%	val loss 0.427, val acc 88.315%
5	epoch 5	train loss 0.164, train acc 95.245%	val loss 0.407, val acc 89.040%
6	epoch 6	train loss 0.137, train acc 96.055%	val loss 0.392, val acc 89.341%
7	epoch 7	train loss 0.117, train acc 96.621%	val loss 0.384, val acc 89.499%
8	epoch 8	train loss 0.101, train acc 97.105%	val loss 0.380, val acc 89.609%
9	epoch 9	train loss 0.088, train acc 97.505%	val loss 0.382, val acc 89.803%
10	epoch 10	train loss 0.077, train acc 97.818%	val loss 0.380, val acc 89.892%
11			
12	test loss 0.392, test acc 89.356%		

Figure 13 shows training and validation loss for 10 epochs. In this model I used pretrained token embedding (GloVe) as from the suggestion in the assignment description, which help the model has better performance. I have tried 100, 200, 300 dimensions of this embedding. They all yielded the similar results. I use Adam as an optimizer. The learning rate that I had tried was 0.01, 0.001, and 0.0001. The best validation accuracy is when learning rate = 0.001. The next hyper-parameter that I had tried was weight_decay. The model yielded very similar result for the values 0.00001, 0.000001, or




Figure 13: Training and validation loss

0. The hidden dimension was the next things that I experimented. The large hidden dimension (64, 128: similar results) gave the higher accuracy rate while the small dimension (10) gave the lower accuracy rate. I tried to run many size of epoch (10, 20, 30). I could not see much different from these 3 values, so I chose the smallest one.

I also kept track of the lowest validation loss model to use in the prediction of test data. This is the early-stopping technique that I used for this work.

The above result and graph's parameter is $lr = 0.001$, $weight\ decay = 0.000001$, $epoch = 10$, $hidden\ dimension = 128$, $embedding\ dimension = 100$, $batch\ size = 128$

TASK 3.3

The results from `tag_sentence` function is shown as below.

1	TAG	TOKEN
2		
3	DET	the
4	ADJ	old
5	NOUN	man
6	DET	the
7	NOUN	boat
8	PUNCT	.
9	-----	
10	TAG	TOKEN
11		
12	PROPN	The
13	ADJ	complex
14	NOUN	houses
15	VERB	married
16	CCONJ	and
17	ADJ	single
18	NOUN	soldiers
19	CCONJ	and
20	PRON	their
21	NOUN	families
22	PUNCT	.
23	-----	

24	TAG	TOKEN
25		
26	PROPN	The
27	NOUN	man
28	PRON	who
29	VERB	hunts
30	VERB	ducks
31	ADP	out
32	ADP	on
33	NOUN	weekends
34	PUNCT	.